



PSL



Hors d'usage vs Canon

Textual Evidence from 11,000 French Novels

Jean Barré

Mapping the Canon: Quantitative Approaches to Literary History

Technical University of Darmstadt — June 19, 2026

LaTTiCe (CNRS, ENS-PSL, Sorbonne Nouvelle)

A small piece of good news

PhD defended — **5 June 2026**

Modeling Genre Literature: Computational Methods and Genre Evolution in French-Language Fiction

Building the corpus

How to assemble a historically valid corpus of genre fiction when metadata is incomplete?

Birth & death of genres

When does a genre become formally recognizable — and when does it fade?

A small piece of good news

PhD defended — **5 June 2026**

Modeling Genre Literature: Computational Methods and Genre Evolution in French-Language Fiction

Building the corpus

How to assemble a historically valid corpus of genre fiction when metadata is incomplete?

Birth & death of genres

When does a genre become formally recognizable — and when does it fade?

*The thesis was about **genre**, not canon — but today I get to bring the two together.*

Cohen (1999): literature “*hors d’usage*”

The serial, mass-market fiction — feuilleton, adventure, detective — that **dominated the market** and then fell **out of literary memory**.

Cohen (1999): literature “*hors d’usage*”

The serial, mass-market fiction — feuilleton, adventure, detective — that **dominated the market** and then fell **out of literary memory**.

Moretti (2000): canon vs archive

Of the thousands of novels published, criticism and the school keep only a **handful**. The rest is the **archive** — read in its time, forgotten since.

Cohen (1999): literature “*hors d’usage*”

The serial, mass-market fiction — feuilleton, adventure, detective — that **dominated the market** and then fell **out of literary memory**.

Moretti (2000): canon vs archive

Of the thousands of novels published, criticism and the school keep only a **handful**. The rest is the **archive** — read in its time, forgotten since.

This talk

Quantitative formalism

What organizes the **hors d’usage** as a textual space?

We want to test two forces — **canon** and **genre** — and ask what the texts themselves reveal.

A paradox: canonical authors of non-canonical genres



Alexandre Dumas

historical romance

Les Trois Mousquetaires
(1844)



Eugène Sue

roman-feuilleton


Les Mystères de Paris (1842)



Jules Verne

scientific adventure

Voyage au centre de la Terre
(1864)



Maurice Leblanc

detective fiction

Arsène Lupin (1905–)

A paradox: canonical authors of non-canonical genres



Alexandre Dumas

historical romance

Les Trois Mousquetaires
(1844)



Eugène Sue

roman-feuilleton


Les Mystères de Paris (1842)



Jules Verne

scientific adventure

Voyage au centre de la Terre
(1864)



Maurice Leblanc

detective fiction

Arsène Lupin (1905–)

All four emerged from **serial, genre-bound, mass-market fiction** —
yet all four are durably **canonical**: still read, adapted, translated.

A paradox: canonical authors of non-canonical genres



Alexandre Dumas

historical romance

Les Trois Mousquetaires
(1844)



Eugène Sue

roman-feuilleton


Les Mystères de Paris (1842)



Jules Verne

scientific adventure

Voyage au centre de la Terre
(1864)



Maurice Leblanc

detective fiction

Arsène Lupin (1905–)

All four emerged from **serial, genre-bound, mass-market fiction** —
yet all four are durably **canonical**: still read, adapted, translated.


Question: does canonicity leave a **measurable textual trace** —
and how does it compare to the trace left by **genre** itself?

What we already know: canon *is* textually measurable


Canon/non-canon differences have been measured quantitatively across **English, French, Dutch, German, Danish...**

What we already know: canon *is* textually measurable

Canon/non-canon differences have been measured quantitatively across **English, French, Dutch, German, Danish...**




English poetry, 1820–1919: prestige is predictable from diction, and stable for a century (Underwood & Sellers 2016).




French novels: probability of being predicted canonical rises over time — production **converges toward the canonical norm** (Barré et al. 2023).

What we already know: canon *is* textually measurable



English poetry, 1820–1919: prestige is predictable from diction, and stable for a century (Underwood & Sellers 2016).




French novels: probability of being predicted canonical rises over time — production **converges toward the canonical norm** (Barré et al. 2023).

This talk's hypothesis

One major factor is **genre** — the *hors d'usage*, the serial popular fiction that makes up most of the archive.

The corpus: the archive, made accessible




Temporal distribution; production peaks in the 1880s (golden age of the *roman-feuilleton*).

“Hors d’usage” corpus

10 898 French novels, 1800–1950

- OCR-filtered (>95% correct words)
- deduplicated (MinHash + metadata)
- BnF / Wikidata author identities

The corpus: the archive, made accessible



Temporal distribution; production peaks in the 1880s (golden age of the *roman-feuilleton*).

“Hors d’usage” corpus

10 898 French novels, 1800–1950

- OCR-filtered (>95% correct words)
- deduplicated (MinHash + metadata)
- BnF / Wikidata author identities


Labels

1428 adventure • **830** detective

~**1160** canonical (~11%)

genre: supervised retrieval; canon: reception proxies
(Barré et al. 2023)

The corpus: the archive, made accessible



Temporal distribution; production peaks in the 1880s (golden age of the *roman-feuilleton*).

“Hors d’usage” corpus

10 898 French novels, 1800–1950

- OCR-filtered (>95% correct words)
- deduplicated (MinHash + metadata)
- BnF / Wikidata author identities

Labels

1428 adventure • **830** detective

~**1160** canonical (~11%)

genre: supervised retrieval; canon: reception proxies
(Barré et al. 2023)

*Openly released — corpus, labels, pipeline — on
Zenodo (JOHD 2026).*

Two candidate forces

Forward similarity = how much a novel resembles the production published *after* it.

What raises it: canonical status, or genre membership?

Two candidate forces

Forward similarity = how much a novel resembles the production published *after* it.
What raises it: canonical status, or genre membership?

Canon

The **author / work** consecrated by
critics, schools, anthologies.

Notion of **Cultural Grammar** (Altieri, 1983)

An institution of *selection*.

Two candidate forces

Forward similarity = how much a novel resembles the production published *after* it.
What raises it: canonical status, or genre membership?

Canon

The **author / work** consecrated by
critics, schools, anthologies.
Notion of **Cultural Grammar** (Altieri, 1983)
An institution of *selection*.

Genre

The **serial formula** — shared scenes,
figures, conventions.
Notion of **Factory fiction** (Letourneux,
2017)
An institution of *production*.

Two candidate forces

Forward similarity = how much a novel resembles the production published *after* it.
What raises it: canonical status, or genre membership?

Canon

The **author / work** consecrated by
critics, schools, anthologies.
Notion of **Cultural Grammar** (Altieri, 1983)
An institution of *selection*.

Genre

The **serial formula** — shared scenes,
figures, conventions.
Notion of **Factory fiction** (Letourneux,
2017)
An institution of *production*.

We measure forward similarity, conditioned on each, against the corpus baseline.

Measuring similarity over the long run

How do we measure how much novels resemble each other across time?

1

Encoder model **fine-tuned on French novels**

Measuring similarity over the long run

How do we measure how much novels resemble each other across time?

1

Encoder model **fine-tuned on French novels**




2

Inference on **all passages** of all novels


Measuring similarity over the long run

How do we measure how much novels resemble each other across time?




Measuring similarity over the long run

How do we measure how much novels resemble each other across time?



Measuring similarity over the long run

How do we measure how much novels resemble each other across time?




A single **novel** $\times$ **novel** similarity matrix — then we read it through **canon**, **non-canon** and **genre** tags.

Result 1 — the “style of the time”, forward

Forward similarity

- For a novel in year y_0 : mean similarity to novels published at $y_0 + t$, $t = 0 \dots 50$
- Forward only — the production that comes *after*
- 5-year rolling mean, ± 1 SE band

Height = closeness to the surrounding production; **slope** = how fast a novel is “left behind”. Read against the grey corpus baseline.




Forward similarity, all novels (5-year rolling mean $\pm$ SE)

Result 1 — the “style of the time”, forward

Forward similarity

- For a novel in year y_0 : mean similarity to novels published at $y_0 + t$, $t = 0 \dots 50$
- Forward only — the production that comes *after*
- 5-year rolling mean, ± 1 SE band


Height = closeness to the surrounding production; **slope** = how fast a novel is “left behind”. Read against the grey corpus baseline.



Forward similarity, all novels (5-year rolling mean $\pm$ SE)


A clear **decay**: novels are most similar to their immediate future, steadily less so over 50 years.

Result 2 — canon is only a faint overlay



Whole corpus: canon vs archive

Result 2 — canon is only a faint overlay




Whole corpus: canon vs archive

Whole corpus

Canon sits **slightly below** archive — a faint signal in the expected direction: canon is marginally *less* absorbed into the later production.

Result 2 — canon is only a faint overlay




Whole corpus: canon vs archive

Whole corpus

Canon sits **slightly below** archive — a faint signal in the expected direction: canon is marginally *less* absorbed into the later production.


The gap is **tiny**, and the bands nearly overlap.

Result 3 — inside a genre: adventure



Within adventure: canon vs archive (vs general literature, grey)

Result 3 — inside a genre: adventure




Within adventure: canon vs archive (vs general literature, grey)

Genre sets the level

Adventure sits **far above** general literature (~ 0.57 vs ~ 0.51): a homogeneous, formula-driven space.

Result 3 — inside a genre: adventure



Within adventure: canon vs archive (vs general literature, grey)


Genre sets the level

Adventure sits **far above** general literature (~ 0.57 vs ~ 0.51): a homogeneous, formula-driven space.

Canon barely shifts it


Same faint ordering *inside* the genre: canon just below archive. The **genre** sets the level; canon is a second-order effect.

Result 4 — inside a genre: detective



Within detective: archive vs (sparse) canon

Result 4 — inside a genre: detective




Within detective: archive vs (sparse) canon

Steepest decay

Detective fiction **crystallizes fast**, then its early texts are quickly left behind.

Result 4 — inside a genre: detective



Within detective: archive vs (sparse) canon


Steepest decay

Detective fiction **crystallizes fast**, then its early texts are quickly left behind.


Canon too sparse

Only **27** canonical detective novels — no stable canon curve. But the **genre signal** still sits far above the baseline.

Even the founders are pulled in: Dumas & Sue




Dumas, *Les Trois Mousquetaires* (1844)




Sue, *Les Mystères de Paris* (1842)

Even the founders are pulled in: Dumas & Sue




Dumas, *Les Trois Mousquetaires* (1844)




Sue, *Les Mystères de Paris* (1842)

The founders of the **roman-feuilleton**: their curves track the corpus, not above it — the *feuilleton* mode they helped invent **became** the mainstream they are measured against.

Two trajectories out of the feuilleton: Verne & Leroux




Verne, *Vingt mille lieues sous les mers* (1871)




Leroux, *Chambre jaune* (1907)

Two trajectories out of the feuilleton: Verne & Leroux



Verne, *Vingt mille lieues sous les mers* (1871)



Leroux, *Chambre jaune* (1907)

Verne anchors a space that detaches into science fiction; **Leroux's** *rising* tail is the whodunit forming *around* him — the closest thing to genre ancestry, but it is the **genre** converging, not the author broadcasting.

What organizes the *hors d'usage*

Within the popular-fiction archive, the dominant textual force is **generic**, not canonical.

Within the popular-fiction archive, the dominant textual force is **generic**, not canonical.

Genre: the organizing principle

Genre **structures** the archive — pulling novels into a shared formal space (Cawelti's *formulas*, Eco's intertextual scenarios, Letourneux's *factory fiction*).

A large, measurable textual effect.

What organizes the *hors d'usage*

Within the popular-fiction archive, the dominant textual force is **generic**, not canonical.

Genre: the organizing principle

Genre **structures** the archive — pulling novels into a shared formal space (Cawelti's *formulas*, Eco's intertextual scenarios, Letourneux's *factory fiction*).

A large, measurable textual effect.

Canon: a faint overlay

Canonicity as a *cultural grammar* — but a **weak textual signal**, nearly orthogonal to the genre's formal gravity.

What organizes the *hors d'usage*

Within the popular-fiction archive, the dominant textual force is **generic**, not canonical.

Genre: the organizing principle

Genre **structures** the archive — pulling novels into a shared formal space (Cawelti's *formulas*, Eco's intertextual scenarios, Letourneux's *factory fiction*).

A large, measurable textual effect.

Canon: a faint overlay

Canonicity as a *cultural grammar* — but a **weak textual signal**, nearly orthogonal to the genre's formal gravity.

Canonization does not follow the textual pull of the genre system — it runs **alongside** it, on other criteria.

Limits

- Novel-mean embeddings capture lexical and thematic dimensions
- Similarity $\neq$ influence: a **family-resemblance** signal
- Right-censoring past +30y (widening bands)

Limits

- Novel-mean embeddings capture lexical and thematic dimensions
- Similarity $\neq$ influence: a **family-resemblance** signal
- Right-censoring past +30y (widening bands)

Perspectives

- Move **beyond the mean**: compare the *full* set of passage vectors, not one average per novel
 $\hookrightarrow$ information theory offers tools to compare **distributions**
- **Between** genres, not just within: does the feuilleton feed the detective novel?

Thank you

Discussion

jean.barre@ens.psl.eu • <https://jeanbarre.eu>

Corpus, labels & code: <https://doi.org/10.5281/zenodo.20233179>