

In Search of Lost Adventure Novels: Supervised Genre Retrieval and Corpus Refinement in Gallica

Task Overview

- Goal: retrieve a "lost" popular genre from a massive digital archive.
- Problem: genre metadata are unreliable; boundaries are porous.
- Corpus: *Hors d'usage: Curated French Fiction from Gallica, 11K novels*
- Evaluation: internal 5-fold CV + external benchmark

Contributions

- ✓ Perspectivist seed list: 101 novels cited by scholars
- ✓ Two-stage pipeline: supervised classification + k-NN graph
- ✓ 1,391-novel adventure corpus released (Zenodo)


Training Data & Cross-Validation

CLASSIFIER	REP.	BACC	F1	AUC
SVM linear	BoW	0.87	0.87	0.94
Random Forest	BoW	0.88	0.88	0.95
HistGradBoost	BoW	0.85	0.85	0.94
SVM linear	Emb*	0.86	0.85	0.93
Random Forest	Emb	0.88	0.88	0.95
HistGradBoost	Emb	0.89	0.89	0.95**

* Emb = mean paragraph embeddings, BGE-M3 fine-tuned on French literary texts (1,024-d)

** Selected for deployment on external F1 (Chapitres) — see Final Results box


Pipeline Architecture



Corpus

Gallica novels (raw)	19,240
After filters	10,998
(OCR>95%, dedup)	
Seed positives	101
(Letourneux)	
Final adventure corpus	1,391


The Corpus as a Graph




Gains from Graph Refinement

Metric	Before	After
Novels	1,375	1,391
Mean similarity	0.789	0.825
Std. deviation	0.098	0.080
Homophily	0.754	0.795

Final Results



Temporal Distribution of the Retrieved Corpus



Perspectives

- Scene-level modeling (danger/exploration) beyond global representations
- Multi-label / segment-granularity genre membership
- Extension to other archive-genres (sentimental, crime fiction)